

Maximum Likelihood Estimation and Confidence Intervals

The Poisson distribution is a discrete probability distribution over the non-negative integers with a single parameter $\lambda > 0$ and PMF:

$$\Pr[X = k] = \frac{\lambda^k e^{-\lambda}}{k!}$$

This distribution is useful in modeling the number of event occurrences in a fixed time interval when the probability of an event occurrence does not depend on the time since the last event. The parameter of this distribution, λ , is the expected number of events in one time interval. As a result, the expectation and variance of a $\text{Pois}(\lambda)$ random variable are both λ .

Suppose that you have a set of samples $D = \{x_1, \dots, x_n\}$. We propose a probability model in which these samples are taken from X_1, \dots, X_n , which are independent $\text{Pois}(\lambda)$ random variables.

1. Write the **log-likelihood** function of λ under D . Use natural log here (base e).

The log-likelihood function :

To write the log likelihood function of λ for a set of i.i.d samples $D = \{x_1, x_2, \dots, x_n\}$ drawn from a poisson distribution, we first need to define the likelihood function and then take its natural logarithm.

H_0 defines the probability of observing data D given a parameter vector θ . for i.i.d samples the likelihood function $L(\theta; D)$ is the product of the probabilities of observing each individual sample:

$$L(\theta; D) = \prod_{i=1}^n p_{\theta}(x_i)$$

So in this case, our $\theta = \lambda$

So the PMF for a poisson distribution is given as:

$$p_{\lambda}(x_i) = \Pr[X = x_i] = (\lambda^{x_i} e^{-\lambda}) / (x_i!)$$

So the likelihood function for the given sample D (under poisson model) is:

$$L(\lambda; D) = \prod_{i=1}^n \frac{\lambda^{x_i} e^{-\lambda}}{x_i!}$$

The log-likelihood function, $l(\lambda; D)$ is obtained by taking the natural logarithm of the likelihood function.

Hence:

$$\begin{aligned} l(\lambda; D) &= \log\left(\prod_{i=1}^n \frac{\lambda^{x_i} e^{-\lambda}}{x_i!}\right) = \sum_{i=1}^n \log\left(\frac{\lambda^{x_i} e^{-\lambda}}{x_i!}\right) = \sum_{i=1}^n \frac{\log(\lambda^{x_i} e^{-\lambda})}{\log(x_i!)} = \sum_{i=1}^n \log(\lambda^{x_i}) + \log(e^{-\lambda}) - \log(x_i!) = \\ &= \sum_{i=1}^n \log(\lambda^{x_i}) - \lambda - \log(x_i!) = \sum_{i=1}^n \log(\lambda^{x_i}) - \sum_{i=1}^n \lambda - \sum_{i=1}^n \log(x_i!) = \\ &= \left(\sum_{i=1}^n x_i\right) \log(\lambda) - n\lambda - \sum_{i=1}^n \log(x_i!) \end{aligned}$$

And as we were asked to use natural logarithm:

$$\left(\sum_{i=1}^n x_i\right) \ln(\lambda) - n\lambda - \sum_{i=1}^n \ln(x_i!)$$

2. Use the **log-likelihood** function you derived above to find an expression for the maximum likelihood estimator (MLE) $\hat{\lambda}$ of λ .

To find the MLE $\hat{\lambda}$ we can take the derivative of the log-likelihood function with respect to λ , set it to 0 and solve for λ . and this way we will find the $\lambda = \theta$ that maximizes the function.

The log-likelihood function of λ for a set of samples $D = \{x_1, x_2, \dots, x_n\}$ from a poisson

distribution was derived as $\left(\sum_{i=1}^n x_i\right) \ln(\lambda) - n\lambda - \sum_{i=1}^n \ln(x_i!)$.

$$\frac{\partial l}{\partial \lambda} = \frac{\partial}{\partial \lambda} \left(\sum_{i=1}^n x_i \ln(\lambda) - n\lambda - \sum_{i=1}^n \ln(x_i!) \right) = 0$$

So now we'll differentiate each term:

$$\frac{\partial}{\partial \lambda} \left(\sum_{i=1}^n x_i \ln(\lambda) \right) = \left(\sum_{i=1}^n x_i \right) \frac{1}{\lambda}$$

$$\frac{\partial}{\partial \lambda} - n\lambda = -n$$

$$\frac{\partial}{\partial \lambda} \sum_{i=1}^n \ln(x_i!) = 0$$

Therefore we get:

$$\frac{\partial}{\partial \lambda} l(\lambda; D) = \left(\sum_{i=1}^n x_i \right) \left(\frac{1}{\lambda} \right) - n$$

$$\text{So: } 0 = \left(\sum_{i=1}^n x_i \right) \left(\frac{1}{\lambda} \right) - n \rightarrow n = \left(\sum_{i=1}^n x_i \right) \left(\frac{1}{\lambda} \right) \rightarrow \frac{n}{\sum_{i=1}^n x_i} = \frac{1}{\lambda} \rightarrow \lambda = \frac{\sum_{i=1}^n x_i}{n}$$

So we get that $\hat{\lambda} = \frac{\sum_{i=1}^n x_i}{n} = \frac{1}{n} \sum_{i=1}^n x_i$

We now wish to find an expression for an interval containing λ with probability at least 0.95 under the model we proposed. Namely, a 95% confidence interval (CI) for λ (around the estimated $\hat{\lambda}$).

3. Consider $\hat{\lambda}$ as a random variable whose randomness stems from the randomness in X_1, \dots, X_n . Write an expression for the mean $\mathbb{E}[\hat{\lambda}]$ and the variance $\text{Var}[\hat{\lambda}]$ as a function of n and λ .

We need to find an expression for an interval containing λ with 95% CI. that means we need to construct a 95% confidence interval for λ around the estimated $\hat{\lambda}$.

For a discrete random variable X , the expected value $E(x)$ is defined as the sum of each possible value multiplied by its probability. A key property of the Poisson distribution is that its mean is equal to λ . so for every $1 \leq i \leq n$, $E(x_i) = \lambda$

$$\text{So we now need to compute } E(\hat{\lambda}) = E\left[\frac{1}{n} \sum_{i=1}^n x_i\right] \rightarrow \frac{1}{n} \sum_{i=1}^n E(x_i)$$

$$\text{So we sub } E(x_i) = \lambda$$

$$= E(\hat{\lambda}) = \frac{1}{n} \sum_{i=1}^n \lambda = \frac{1}{n} n\lambda = \lambda$$

So now we find the variance:

The variance of a random variable X is defined as $\text{Var}(X) = E[(X - E[X])^2]$

Since X_1, \dots, X_n are independent $\text{Pois}(\lambda)$ random variables, each X_i has a variance $\text{Var}[X_i] = \lambda$.

$$\text{So we get: } \text{Var}[\hat{\lambda}] = \text{Var}\left[\frac{1}{n} \sum_{i=1}^n x_i\right]$$

We use $\text{Var}[cX] = c^2 \text{Var}[X]$ and the fact that for i.i.d variables we get $\text{Var}[\sum X_i] = \sum \text{Var}[X_i]$

So we get :

$$\text{Var}[\hat{\lambda}] = \left(\frac{1}{n}\right)^2 \text{Var}\left[\sum_{i=1}^n x_i\right] = \left(\frac{1}{n^2}\right) \sum_{i=1}^n \text{Var}[x_i] - (\text{Var}[x_i] = \lambda) \rightarrow \left(\frac{1}{n^2}\right) \sum_{i=1}^n \lambda = \frac{1}{n^2} n\lambda = \frac{\lambda}{n}$$

So $\text{Var}[\hat{\lambda}] = \frac{\lambda}{n}$ and $E(\hat{\lambda}) = \lambda$

4. Assume that your estimator $\hat{\lambda}$ is asymptotically normal, meaning that its **standardized** value approaches a normal distribution as the number of data points goes to infinity:

$$\frac{\hat{\lambda} - \mathbb{E}[\hat{\lambda}]}{\sqrt{\text{Var}[\hat{\lambda}]}} \sim \mathcal{N}(0, 1). \quad (1)$$

(Note: this is implied by the central limit theorem.)

Under this assumption, write expressions for the lower and upper boundaries of the 95% confidence interval for λ around the estimated $\hat{\lambda}$. Your two expressions should be specified as functions of $\hat{\lambda}$, n , and the inverse Gaussian CDF $\Phi^{-1}(p)$.

Given that the standardized value of $\hat{\lambda}$ follows a standard normal distribution for large n :

$$Z = \frac{(\hat{\lambda} - E(\hat{\lambda}))}{\sqrt{\text{Var}[\hat{\lambda}]}} \sim N(0, 1)$$

So given our mean and variance we get $Z = \frac{\hat{\lambda} - \lambda}{\sqrt{\frac{\lambda}{n}}}$

So now we need to find for which z we get $P(-z \leq Z \leq z) = 0.95$

So we get $P(-z \leq Z \leq z) = \Phi(z) - (1 - \Phi(z)) = 2\Phi(z) - 1$

So setting $2\Phi(z) - 1 = 0.95 \rightarrow 2\Phi(z) = 1.95 \rightarrow \Phi(z) = 0.975$

So we've seen in class :

Compute z for confidence level $1 - \alpha = 95\%$:

$$z = \Phi^{-1}\left(1 - \frac{\alpha}{2}\right) = \Phi^{-1}(0.975) \approx 1.96$$

So:

$$-\Phi^{-1}(0.975) \leq \left(\frac{\hat{\lambda} - \lambda}{\sqrt{\frac{\lambda}{n}}}\right) \leq \Phi^{-1}(0.975) \rightarrow -\Phi^{-1}(0.975)\left(\sqrt{\frac{\hat{\lambda}}{n}}\right) \leq \hat{\lambda} - \lambda \leq \Phi^{-1}(0.975)\left(\sqrt{\frac{\hat{\lambda}}{n}}\right) \rightarrow$$

$$-\Phi^{-1}(0.975)\left(\sqrt{\frac{\hat{\lambda}}{n}}\right) - \hat{\lambda} \leq \lambda \leq \Phi^{-1}(0.975)\left(\sqrt{\frac{\hat{\lambda}}{n}}\right) - \hat{\lambda} \rightarrow$$

$$= \hat{\lambda} - \Phi^{-1}(0.975)\left(\sqrt{\frac{\hat{\lambda}}{n}}\right) \leq \lambda \leq \hat{\lambda} + \Phi^{-1}(0.975)\left(\sqrt{\frac{\hat{\lambda}}{n}}\right)$$

So we get that the interval is:

$$[\hat{\lambda} - \phi^{-1}(0.975)(\frac{\sqrt{\hat{\lambda}}}{n}), \hat{\lambda} + \phi^{-1}(0.975)(\frac{\sqrt{\hat{\lambda}}}{n})]=$$

$$[\hat{\lambda} - 1.96(\frac{\sqrt{\hat{\lambda}}}{n}), \hat{\lambda} + 1.96(\frac{\sqrt{\hat{\lambda}}}{n})]$$